

MEIC 2020/2021  
Aprendizagem - Machine Learning  
Homework II  
*Prof. Andreas Wichert*  
Deadline 03/05/2021  
*Submit on Fenix as pdf*

Submission instructions:

Submit a single .zip file on Fenix containing:

- 1 pdf with a 2 page report
- 1 with the working code for Exercise 2

Every file name must be the group number

You can install all the libraries using <https://www.anaconda.com>, or you can use google colab, <https://colab.research.google.com/notebooks/intro.ipynb>

Exercise 1 (2 pts)

Use the famous dataset that contains 150 iris flowers to detect Iris-Virginica or Not Iris-Virginica (**two classes**). Perform Logistic Regression with Scikit-learn using the code snippet provided below. Indicate the training and test accuracy (fraction of correctly classified examples).

Example of code using NumPy and Scikit-learn (<https://scikit-learn.org/stable/>):

```
from sklearn import datasets
from sklearn.model_selection import train_test_split
import numpy as np

iris = datasets.load_iris()
print("Iris Data\n", list(iris.keys()))

X = iris.data[:, :2] # we only take the first two features.

y = (iris["target"] == 2).astype(np.int) # 1 if Iris-Virginica, else 0

print(y)

# Split into a train and test partition.
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
random_state=0)

from sklearn.linear_model import LogisticRegression

log_reg = LogisticRegression()

log_reg.fit(X_train, y_train)

print("Predict (train):", log_reg.predict(X_train[:len(y_train), :]))
print("Predict (test):", log_reg.predict(X_test[:len(y_test), :]))
```

MEIC 2020/2021  
Aprendizagem - Machine Learning  
Homework II  
*Prof. Andreas Wichert*  
Deadline 03/05/2021  
*Submit on Fenix as pdf*

Exercise 2 (18 pts)

Use the same Iris dataset (**two classes**) to investigate **gradient descent** of one unit (perceptron). Implement the perceptron in Python and NumPy (without Scikit-learn). You can still use the code snippet of the previous exercise **but only to load and split the data**. You will try three configurations:

- i. Squared error and linear unit (Linear Regression)
- ii. Squared error and sigmoid unit
- iii. Cross-entropy error and sigmoid unit (Logistic Regression)

$$\sigma(net) = \frac{1}{1 + e^{(-net)}} = \frac{e^{(net)}}{1 + e^{(net)}}$$

The sigmoid (logistic) activation function.

- a) (3 pts) Describe the learning rule for each example i, ii, iii, implement it using Python and NumPy (not Scikit-learn).
- b) (4 pts) For each model perform three experiments with different random initialization of weights with a fixed learning rate (does it have any consequences?)
- c) (4 pts) Which of the three models i, ii, iii converge faster to a near-optimal solution? Indicate the number of steps/epochs related to the learning rate (three experiments for each model). For the best choice of learning rate for each model, plot the error curves as a function of the number of epochs. You can use matplotlib for plotting, to do it store the data into an array A:

```
import matplotlib.pyplot as plt
```

```
plt.plot(A)
```

```
plt.show()
```

- d) (4 pts) Which model is the most accurate? Report the training and test accuracies. Are the models equivalent or not and why do you think so?
- e) (3 pts) Do you get the same results as Logistic Regression of Scikit-learn (Exercise 1)?